

A Dynamic Dictionary for Discovering Indirect Translation Equivalents

Bogdan Babych, Anthony Hartley, Serge Sharoff

Centre for Translation Studies
University of Leeds, UK
{b.babych,a.hartley,s.sharoff}@leeds.ac.uk

Abstract

We present the design and evaluation of a novel software application intended to help translators with rendering problematic expressions from the general lexicon. It does this dynamically by first generalising the problem expression in the source language and then searching for possible translations in a large comparable corpus. These candidate solutions are ranked and presented to the user. The method relies on measures of distributional similarity and on bilingual dictionaries. It outperforms established techniques for extracting translation equivalents from parallel corpora. The interface to the system is available at: <http://corpus.leeds.ac.uk/assist/v05/>.

1 Introduction

This paper describes ASSIST, a system designed to assist humans in translating expressions that do not necessarily have a word-for-word (‘compositional’) equivalent in the target language (TL). In the spirit of (Kay, 1997), it is intended as a *translator’s amenuensis*, ‘under the tight control of a human translator ... to help increase his productivity and not to supplant him’.

One area where human translators particularly appreciate assistance is in the translation of expressions from the general lexicon. Unlike technical terms – which generally share the same part-of-speech (POS) across languages and, in the ideal case, respect the rule ‘one form one meaning’ – the contextually appropriate equivalents of general language expressions are often indirect and open to variation. Moreover, human translators, even in non-literary fields, often value legitimate variation. Thus the French expression *il faillit échouer* (lit.: he faltered to fail) may be variously rendered as *he almost/nearly/all but failed*; *he was on the verge/brink of failing/failure*; *failure loomed*. All of these translations are indirect in that they involve lexical shifts or POS transformations.

Finding such translations is a hard task that can benefit from automated assistance. ‘Mining’ such indirect equivalents is difficult, precisely because of the structural mismatch, but also because of the paucity of suitable aligned (parallel) corpora. The approach adopted here includes the use of *comparable* corpora in source and target languages, i.e. corpora of texts dealing with similar subject matter and intended for similar readerships. These are relatively easy to create, by ‘harvesting them from the internet, for example, and there are no alignment costs. The greatest challenge is to generate a list of solutions that translators will find usable and to rank them such that the best are at the top.

While ASSIST is unlike statistical machine translation (SMT – Och and Ney, 2003), where lexical selection is effected by a translation model based on aligned, parallel corpora, the novel techniques it has developed are exploitable in the SMT paradigm. It also differs from

now traditional uses of comparable corpora for detecting translation equivalents (Rapp, 1999) or extracting terms (Grefenstette, 2002) which exhibit a one-to-one correspondence irrespective of the context. ASSIST addresses difficulties with expressions from the general lexicon, whose translation is context-dependent.

2 Methodology

The software acts as a decision support system for translators. It integrates different technologies for extracting indirect translation equivalents from large comparable corpora. In the following subsections we give the user perspective on the system and describe the methodology underlying each of its sub-tasks. Explanations of some of the technical details are provided by (Babych et al., 2007).

1.1 User perspective

Unlike traditional dictionaries, the system is a *dynamic translation resource* in that it can successfully find translation equivalents for units which have not been stored in advance, even for idiosyncratic multiword expressions which almost certainly will not figure in a dictionary. While ASSIST can rectify gaps and omissions in static lexicographical resources, its major advantage is that it is able to cope with an open set of translation problems, searching for translation equivalents in comparable corpora in runtime. This makes it more than just an extended dictionary.

Contextual descriptors

From the user's perspective the system extracts indirect translation equivalents as sets of *contextual descriptors* – content words that are lexically central in a given sentence, phrase or construction. The choice of these descriptors may determine the general syntactic perspective of the sentence and the use of supporting lexical items. Many translation problems arise from the fact that the mapping between such descriptors across languages is not straightforward, as our earlier French-English example showed.

The system is designed to find possible indirect mappings between sets of descriptors and to verify the acceptability of the mapping into the TL. For example, in the following Russian sentence, the bolded contextual descriptors require indirect translation into English.

*Дети посещают **плохо отремонтированные** школы, в которых **недостает** самого **необходимого***

*(Children attend **badly repaired** schools, in which [it] is **missing** the most **necessary**)*

Combining direct translation equivalents of these words (e.g., translations found in the Oxford Russian Dictionary – ORD) may produce a non-natural English sentence, like the literal translation given above. In such cases human translators usually apply structural and lexical transformations, for instance changing the descriptors' POS and/or replacing them with near-synonyms which fit together in the context of a TL sentence (Munday, 2001: 57-58). Thus, a structural transformation of *плохо отремонтированные* (badly repaired) may give *in poor repair* while a lexical transformation of *недостает самого необходимого* ([it] is missing the most necessary) could give *lacking basic essentials*.

ASSIST can generate such transformations of the descriptors into the TL and checks whether combinations of the translated descriptors actually occur in the TL corpus.

Using the system

Human translators submit queries in the form of one or more SL descriptors which, in their opinion, may require indirect translation. When translators use the system for translating into their native language, the descriptors generated are usually sufficient for them to produce a

correct TL construction or phrase around them (even though the descriptors as presented do not always form a naturally sounding expression). When translators work into a non-native language, they often find it useful to generate concordances for the generated descriptors in order to verify their usage within TL constructions. ASSIST offers these concordances.

For example, for the sentence above translators may submit two queries: *плохо отремонтированные* (badly repaired) and *недостаёт необходимого* (missing necessary). For the first query the system returns a list of descriptor pairs (with information on their frequency in the English corpus) ranked by distributional proximity to the original query, which we explain in Section 1.2. At the top of the list come:

```
bad repair = 30          (11.005)
bad maintenance = 16      (5.301)
bad restoration = 2        (5.079)
poor repair = 60          (5.026)...
```

Underlined hyperlinks lead translators to actual contexts in the English corpus, e.g., *poor repair* generates a concordance containing a desirable TL construction which is a structural transformation of the SL query:

in such a	poor	state of repair
bridge in as	poor	a state of repair as the highways
building in	poor	repair.
dwellings are in	poor	repair;

Similarly, the result of the second query may give the translators ideas for possible lexical transformations:

```
missing need = 14          (5.035)
important missing = 8        (2.930)
missing vital = 8          (2.322)
lack necessary = 204        (1.982)
essential lack = 86         (0.908)...
```

The concordance for the last pair of descriptors contains the phrase *they lack the three essentials*, which illustrates the transformation. The resulting translation may be the following:

*Children attend schools that are in **poor repair** and **lacking** basic **essentials***

Thus ASSIST supports translators in making decisions about indirect translation equivalents in a number of ways: it suggests possible structural and lexical transformations for contextual descriptors; it verifies which translation variants co-occur in the TL corpus; and it illustrates the use of the transformed TL lexical descriptors in actual contexts.

1.2 Generating and ranking translation equivalents

The method for generating translation equivalents is a generalisation of one used in previous work (Sharoff et al., 2006) on extracting equivalents for continuous multiword expressions (MWEs). Essentially, the method expands the search queries for each word and its dictionary translations with entries from thesauri automatically computed from the corpora. It then checks which combinations are possible in the TL corpus or corpora. These potential translation equivalents are now ranked by their distributional similarity to the original query and presented to the user. In this way, the range of retrievable equivalents has been extended from a relatively limited range of two-word constructions which mirror POS categories in SL and TL to a much wider set of co-occurring lexical content items, which may appear in a different order, at some distance from each other, and belong to different POS categories.

The method works best for expressions from the general lexicon which do not have established equivalents. We have more recently extended it to find terminology. It relies on a high-quality bilingual dictionary (en-ru ~30K words, ru-en ~50K words, combining ORD and the core part of the Multitran online dictionary) and large comparable corpora (~200M words of English, ~70M words of Russian) of news texts.

For each of the SL query terms q the system generates its *dictionary translation* $Tr(q)$ and its *similarity class* $S(q)$ – a set of words with a similar distribution in a monolingual corpus. The descriptor and each word in the similarity class are then translated into the TL using ORD or the Multitran dictionary, resulting in a TL set of descriptors which combines $Tr(q)$ and $Tr(S(q))$. On the TL side we further generate the similarity class $S(Tr(q))$ of the dictionary translations of each query term q , but only for dictionary translations of query terms $Tr(q)$. We refer to the resulting set of TL words as a *translation class* T .

$$T = \{Tr(q) \cup Tr(S(q)) \cup S(Tr(q))\}$$

Translation classes approximate lexical and structural transformations which can potentially be applied to each of the query terms. Automatically computed similarity classes do not require resources like WordNet (<http://wordnet.princeton.edu/>), and they are much more suitable for modelling translation transformations, since they often contain a wider range of words of different POS which share the same context. For example, the similarity class of the word *lack* contains words such as *absence, insufficient, inadequate, lost, shortage, failure, paucity, poor, weakness, inability, need*. This clearly goes beyond the range of traditional thesauri.

For multiword queries, the system performs a consistency check on all theoretically possible combinations of words from the translation classes of each of the SL query words, discarding those combinations which are not found in a database of discontinuous content word bi-grams that actually occur in the TL corpus. The database contains the set of all bi-grams that occur in the corpus with a frequency ≥ 4 within a window of 5 words (over 9M bigrams for each language).

Larger N-grams ($N > 2$) in queries are split into combinations of bi-grams, which we found to be an optimal solution to the problem of the scarcity of higher order N-grams in the corpus. Thus, for the query *gain significant importance* ASSIST considers three word pairs – (*significant importance*), (*gain importance*), (*gain significant*) – which enables it to find an indirect equivalent *получить весомое значение* (lit.: receive weighty meaning).

Despite the consistency checking, the set of potential translation equivalents may still be large and contain much noise. Typically the set contains several hundred elements, of which only a few are really useful for translation.

To make the system usable in practice, i.e., to get useful solutions to appear close to the top (preferably on the first screen of the output), we developed methods of ranking and filtering the returned TL contextual descriptor pairs. The system ranks the returned list of contextual descriptors by their distributional proximity to the original query. Thus, words whose equivalents show similar usage in a comparable corpus receive the highest scores. These scores are computed for each individual word in the output, so there are several ways to combine them to weight words in translation classes and word combinations in the returned list of descriptors.

For example, the system gives the following ranking for the indirect translation equivalents of the Russian phrase *весомое значение* (lit.: weighty meaning) – figures in brackets represent the weighted scores for each pair of TL descriptors:

1. significant importance = 7 (3.610)
2. significant value = 128 (3.211)

- 3. measurable value = 6 (2.657)
- 8. dramatic importance = 2 (2.028)
- 9. important significant = 70 (2.014)
- 10. convincing importance = 6 (1.843)

1.3 Semantic filtering

Ranking of translation candidates can be further improved when translators use an option to filter the returned list by certain lexical criteria, e.g., to display only those examples that contain a certain lexical item, or to require one of the items to be a dictionary translation of the query term. However, lexical filtering is often too restrictive: in many cases translators need to see a number of related words from the same semantic field or subject domain, without knowing the lexical items in advance. In this section we present the semantic filter, which is based on Russian and English semantic taggers which use the same semantic field taxonomy for both languages.

The semantic filter displays only those items which have specified semantic field tags or tag combinations; it can be applied to one or both words in each translation suggestion. The default setting for the semantic filter is the requirement for both words in the resulting TL candidates to contain any of the semantic field tags from a SL query term.

In the next section we present evaluation results for this default setting (which is applied when the user clicks the *Semantic Filter* button), but human translators have further options – to filter by tags of individual words, to use semantic classes from SL or TL terms, etc.

For example, applying the default semantic filter for the output of the query *плохо отремонтированные* (badly repaired) removes the highlighted items from the list:

- 1. bad repair = 30 (11.005)
- 2. good repair = 154 (8.884)]**
- 3. bad rebuild = 6 (5.920)
- 4. bad maintenance = 16 (5.301)]**
- 5. bad restoration = 2 (5.079)
- 6. poor repair = 60 (5.026)
- 7. good rebuild = 38 (4.779)]**
- 8. bad construction = 14 (4.779)

Items 2 and 7 are generated by the system because *good*, *well* and *bad* are in the same similarity cluster for many words (they often share the same collocations). The semantic filter removes examples with *good* and *well* on the grounds that they do not have any of the tags which come from the word *плохо* (badly): in particular, instead of tag A5- (Evaluation: Negative) they have tag A5+ (Evaluation: Positive). Item 4 is removed on the grounds that the words *отремонтированный* (repaired) and *maintenance* do not have any tags in common – they appear ontologically too far apart from the point of view of the semantic tagger.

The core of the system's multilingual semantic tagging is a knowledge base in which single words and MWEs are mapped to their potential semantic field categories. Often a lexical item is mapped to multiple semantic categories, reflecting its potential multiple senses. In such cases, the tags are arranged by the order of likelihood of meanings, with the most prominent first.

3 Objective evaluation

In the *objective evaluation* we tested the performance of our system on a selection of indirect translation problems, extracted from a parallel corpus consisting mostly of articles from English and Russian newspapers (118,497 words in the R-E direction, 589,055 words in the E-R direction). It was aligned at the sentence level by JAPA (Langlais et al., 1998), and further at the word level by GIZA++ (Och and Ney, 2003).

1.4 Comparative performance

The intuition behind the objective evaluation experiment is that the capacity of ASSIST to find indirect translation equivalents in comparable corpora can be compared with the results of automatic alignment of parallel texts used in translation models in SMT: one of the major advantages of the SMT paradigm is its ability to reuse indirect equivalents found in parallel corpora (equivalents that may never come up in hand-crafted dictionaries). Thus, automatically generated GIZA++ dictionaries with word alignment contain many examples of indirect translation equivalents.

We used these dictionaries to simulate the generator of translation classes T , which we recombined to construct the set of potential translation equivalents, similarly to the procedure ASSIST uses to generate its output. However, the two approaches generate indirect translation suggestions on the basis of radically different material: the GIZA dictionary uses evidence from parallel corpora of existing human translations, while our system recombines translation candidates on the basis of their distributional similarity in monolingual comparable corpora. Since GIZA represents the current state of the art, we took it as a baseline.

Translation problems for the objective evaluation experiment were manually extracted from two parallel corpora: a section of about 10,000 words of a corpus of English and Russian newspapers, which we also used to train GIZA to find equivalents; and a section of the same length from a corpus of interviews published on the Euronews.net website.

We selected expressions which represented cases of lexical transformations (as illustrated in Section 1), containing at least two content words both in the SL and TL. These expressions were converted into pairs of contextual descriptors – e.g., *recent success*, *reflect success* – and submitted to the system and to the GIZA dictionary. We compared the ability of ASSIST and of GIZA to find indirect translation equivalents which matched the equivalents used by human translators. The output from both systems was checked to see whether it contained the contextual descriptors used by human translators. We submitted 388 pairs of descriptors extracted from the newspaper translation corpus and 174 pairs extracted from the Euronews interview corpus. Half of these pairs were Russian, and the other half English.

We computed recall figures for 2-word combinations of contextual descriptors and single descriptors within those combinations. By recall, we mean the percentage of human-generated descriptors that were proposed by ASSIST and by GIZA. We also show the recall of translation variants provided by the ORD on the same data set. For example, for the query *недостаёт необходимого* ([it] is missing necessary [things]) human translators give the solution *lacking essentials*, the lemmatised descriptors being *lack* and *essential*. ORD returns direct translation equivalents *missing* and *necessary*. The GIZA dictionary additionally contains several translation equivalents for the second term (with alignment probabilities) including: *necessary* ~0.332, *need* ~0.226, *essential* ~0.023. ASSIST returns both descriptors used in human translation as a pair – *lack essential* (ranked 41 without filtering and 22 with the default semantic filter). Thus, for a 2-word combination of the descriptors only the output of our system matched the human solution, which we counted as one hit for ASSIST and no hits for ORD or GIZA. For 1-word descriptors we counted 2 hits for ASSIST (both words in the human solution are matched), and 1 hit for GIZA – it matches the word *essential* ~0.023 (which also illustrates its ability to find indirect translation equivalents).

	2-wd descriptors		1-wd descriptors	
	news	i'view	news	i'view
ORD	6.7%	4.6%	32.9%	29.3%
GIZA++	13.9%	3.4%	35.6%	29.0%
ASSIST	21.9 %	19.5 %	55.8 %	49.4 %

Table 1: Conservative estimate of recall

It can be seen from Table 1 that for the newspaper corpus on which it was trained, GIZA covers a wider set of indirect translation variants than ORD. But ASSIST's recall is even better, both for 2-word and 1-word descriptors.

However, note that GIZA's ability to retrieve from the newspaper corpus certain indirect translation equivalents may be due to the fact that it has previously seen them frequently enough to generate a correct alignment and the corresponding dictionary entry.

The Euronews interview corpus was not used for training GIZA. It represents spoken language and can be expected to contain more 'radical' transformations. The small decline in ORD figures here can be attributed to the fact that there is a difference in genre between written and spoken texts and consequently between transformation types in them. However, the performance of GIZA drops radically on unseen text and becomes approximately the same as ORD.

This shows that indirect translation equivalents in the parallel corpus used for training GIZA are too sparse to be learnt one by one and successfully applied to unseen data, since solutions which fit one context do not necessarily suit others.

The performance of ASSIST stays at about the same level for this new type of text; the decline in its performance is comparable to the decline in ORD figures, and can again be explained by the differences in genre.

1.5 Evaluation of hypothesis ranking

As we mentioned in Section 1.2, correct ranking of translation candidates improves the usability of the system. Again, the objective evaluation experiment gave only a conservative estimate of ranking, because there may be many more useful indirect solutions further up the list in the output of the system which are legitimate variants of the solutions found in the parallel corpus. Therefore, evaluation figures should be interpreted in a comparative rather than an absolute sense.

We use ranking by frequency as a baseline against which to compare the rankings obtained by the method described in Section Error: Reference source not found – by distributional similarity between a candidate and the original query.

Table 2 shows the average rank of human solutions found in parallel corpora and the recall of these solutions for the top 300 examples. Since there are no substantial differences between the figures for the newspaper texts and for the interviews, we report the results jointly for 556 translation problems in both selections (lower rank figures are better).

	Recall	Average rank
2-word descriptors		
frequency (baseline)	16.7 %	<i>rank=93.7</i>
distributional similarity	19.5 %	<i>rank=44.4</i>
similarity + filter	14.4 %	<i>rank=26.7</i>
1-word descriptors		
frequency (baseline)	48.2 %	<i>rank=42.7</i>
distributional similarity	52.8 %	<i>rank=21.6</i>
similarity + filter	44.1 %	<i>rank=11.3</i>

Table 2: Ranking: frequency, similarity and filter

It can be seen from the table that ranking by similarity yields almost a twofold improvement for the average rank figures compared to the baseline. There is also a small improvement in recall, since there is a greater number of relevant examples that appear within the top 300 entries.

The semantic filter once again gives an almost twofold improvement in ranking, since it removes many noisy items. The average is now within the top 30 items, which means that there is a high chance that a translation solution will be displayed on the first screen. The price for improved ranking is decline in recall, since it may remove some relevant lexical transformations if they appear to be ontologically too far apart. But the decline is smaller – about 26.2% for 2-word descriptors and 16.5% for 1-word descriptors. The semantic filter is an optional tool, which can be used to great effect on noisy output: its improvement of ranking outweighs the decline in recall.

Note that the distribution of ranks is not normal, so in Figure 1 we present frequency polygons for rank groups of 30 (which is the number of items that fit on a single screen, i.e., the number of items in the first group (r030) shows solutions that will be displayed on the first screen). The majority of solutions ranked by similarity appear high in the list (in fact, on the first two or three screens).




Figure 1: Polygons for ranks

4 Subjective evaluation

The objective evaluation reported above uses a single human reference translation and is correspondingly conservative in estimating the coverage of the system. However, many expressions studied have more than one fluent translation. For instance, *in poor repair* is not the only equivalent for the Russian expression *плохо отремонтированные*. It is also possible to translate it as *unsatisfactory condition*, *bad state of repair*, *badly in need of repair*, and so on. The objective evaluation shows that the system has been able to find the suggestion used by a particular translator for the problem studied. It does not tell us whether the system has found some other translations suitable for the context. Such legitimate translation variation implies that the performance of a system should be studied on the basis of multiple reference translations. For the purposes of evaluating a fully automatic MT tool, the typical practice of using just two reference translations may be sufficient (Papineni, et al, 2001). However, in the context of a translator's amanuensis which deals with expressions difficult for human translators, it is reasonable to work with a larger range of acceptable target expressions.

With this in mind we evaluated the performance of the tool with a panel of 12 professional translators, members of ITI and the Chartered Institute of Linguists. Test materials were provided in which problematic expressions were highlighted and the translators were asked to find suitable suggestions produced by the tool for these expressions and rank their usability on a scale from 1 to 5 ('not acceptable' to 'fully idiomatic', so '1' means that no usable translation was found at all).

The test sentences themselves were selected from problems discussed on the professional translation forums proz.com and forum.lingvo.ru. Given the range of corpora used in the system (reference and newspaper corpora), the examples were filtered to restrict them to expressions plausible in newspapers.

The goal of the subjective evaluation was to establish the usefulness of the system for translators beyond the conservative estimate given by the objective evaluation. The intuition behind the experiment is that if there are several admissible translations for the SL contextual descriptors, and system output matches any one of these solutions, then the system has generated something useful. Therefore, we computed recall on sets of human solutions rather than on individual solutions. We matched 210 different human solutions to 36 translation problems. To compute more realistic recall figures, we counted cases when the system output matches any of the human solutions in the set. Table 3 compares the conservative estimate of the objective evaluation and the more realistic estimate on a single data set.

	2-wd default		2-wd with sem filter	
	Recall	Av. rank	Recall	Av. rank
Conservative	32.4%	53.68	21.9%	34.67
Realistic	75.0%	7.48	61.1%	3.95

Table 3 Recall and rank for 2-word descriptors

Since the data set is different, the figures for the conservative estimate are higher than those for the objective evaluation data set. However, the table shows there is a gap between the conservative estimate and the realistic coverage of the translation problems by the system, and that real coverage of indirect translation equivalents is potentially much higher.

Table 4 shows averages (and standard deviation σ) of the usability scores (on the scale 1-5) divided into four groups: (1) solutions that are found both by ASSIST and ORD; (2) solutions found only by our system; (3) solutions found only by ORD (4) solutions found by neither:

	ASSIST (+)	ASSIST (-)
ORD (+)	4.03 (0.42)	3.62 (0.89)
ORD (-)	4.25 (0.79)	3.15 (1.15)

Table 4 Human scores and σ for system output

It can be seen from the table that human users find the system most useful for those problems where the solution does not match any of the direct dictionary equivalents, but is generated by the system.

5 Conclusions

We have presented a method of finding indirect translation equivalents in comparable corpora, which has been integrated into a system which assists translators in indirect lexical transfer. The method outperforms established methods of extracting indirect translation equivalents from parallel corpora.

We can interpret these results as an indication that our method, rather than learning individual indirect transformations, models the entire family of transformations entailed by indirect lexical transfer. In other words it learns a translation strategy which is based on the distributional similarity of words in a monolingual corpus, and applies this strategy to novel, previously unseen examples.

The coverage of the tool and additional filtering techniques make it useful for professional translators in automating the search for non-trivial, indirect translation equivalents, especially equivalents for multiword expressions.

Although developed for English and Russian, the ASSIST architecture has been extended to cover English-German translation. This required less than one person-month of effort. Informal evaluations showed the performance to be less good than for English-Russian, but this can be attributed to the inferior quality of the bilingual dictionary used.

More recent work has extended the techniques described here to automatically identifying equivalent single-word and multiword terminological expressions, using large English and French corpora in the CAD-CAM domain.

Acknowledgments

We would like to thank the professional translators who kindly participated in our evaluation trials. This work was supported by EPSRC grant EP/C005902/1 and was conducted jointly with Paul Rayson. Olga Moudraya and Scott Piao of Lancaster University InfoLab.

References

- Bogdan Babych, Anthony Hartley, Serge Sharoff and Olga Moudraya. 2007. Assisting translators in indirect lexical transfer. In *Procs. ACL 2007*, pp. 136-143.
- Gregory Grefenstette. 2002. Multilingual corpus-based extraction and the very large lexicon. In: Lars Borin, editor, *Language and Computers, Parallel corpora, parallel worlds*, pages 137-149. Rodopi.
- Martin Kay. 1997. The proper place of men and machines in language translation. *Machine Translation*, 12(1-2):3-23.
- Philippe Langlais, Michel Simard, and Jean Véronis. 1998. Methods and practical issues in evaluating alignment techniques. In *Procs. Joint COLING-ACL-98*, pp. 711-717.

- Jeremy Munday. 2001. *Introducing translation studies. Theories and Applications*. Routledge, New York.
- Franz Josef Och and Hermann Ney. 2003. A systematic comparison of various statistical alignment models. *Computational Linguistics*, 29(1):19-51.
- Kishore Papineni, Salim Roukos, T. Ward, and W.-J. Zhu. 2001. BLEU: a method for automatic evaluation of machine translation, *RC22176 W0109-022*: IBM.
- Reinhard Rapp. 1999. Automatic identification of word translations from unrelated English and German corpora. In *Procs. ACL-99*, pp. 395-398.
- Reinhard Rapp. 2004. A freely available automatically generated thesaurus of related words. In *Procs. LREC 2004*, pp. 395-398, Lisbon.
- Serge Sharoff, Bogdan Babych and Anthony Hartley 2006. Using Comparable Corpora to Solve Problems Difficult for Human Translators. In *Procs. Joint COLING-ACL-2006 Main Conference Poster Sessions*, pp. 739-746.